

1. I first import my data file using the whole path name.
2. I then use the `info()` to see what types of data type im working with. This also gives me my columns.

```
import pandas as pd

water_pollution = pd.read_csv("/Users/tonydao/Documents/PollutionProject/water_pollution_disease.csv")

water_pollution.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 24 columns):
 #   Column                                                                 Non-Null Count  Dtype  
---  --
 0   Country                                                                3000 non-null   object  
 1   Region                                                                3000 non-null   object  
 2   Year                                                                3000 non-null   int64   
 3   Water Source Type                                                    3000 non-null   object  
 4   Contaminant Level (ppm)                                              3000 non-null   float64  
 5   pH Level                                                            3000 non-null   float64  
 6   Turbidity (NTU)                                                     3000 non-null   float64  
 7   Dissolved Oxygen (mg/L)                                             3000 non-null   float64  
 8   Nitrate Level (mg/L)                                                3000 non-null   float64  
 9   Lead Concentration (µg/L)                                           3000 non-null   float64  
10   Bacteria Count (CFU/mL)                                             3000 non-null   int64   
11   Water Treatment Method                                              3000 non-null   object  
12   Access to Clean Water (% of Population)                            3000 non-null   float64  
13   Diarrheal Cases per 100,000 people                                3000 non-null   int64   
14   Cholera Cases per 100,000 people                                  3000 non-null   int64   
15   Typhoid Cases per 100,000 people                                  3000 non-null   int64   
16   Infant Mortality Rate (per 1,000 live births)                     3000 non-null   float64  
17   GDP per Capita (USD)                                                3000 non-null   int64   
18   Healthcare Access Index (0-100)                                    3000 non-null   float64  
19   Urbanization Rate (%)                                              3000 non-null   float64  
20   Sanitation Coverage (% of Population)                             3000 non-null   float64  
21   Rainfall (mm per year)                                             3000 non-null   int64   
22   Temperature (°C)                                                  3000 non-null   float64  
23   Population Density (people per km²)                               3000 non-null   int64   
dtypes: float64(12), int64(8), object(4)
memory usage: 562.6+ KB
```

3. I notice that Year is a column choice that would be good to start out as a sorted index. This would allow me to conduct basic chart/graph observation when comparing it with other column variables.

```
water_pollution.sort_values("Year", inplace = True)

water_pollution.set_index("Year", inplace= True)
```

- 4: I used the `.head()` and `.tail()` to see if my index has been changed and if it is sorted corrected.

```
[27]: water_pollution.head(5)

[27]:
```

	Country	Region	Water Source Type	Contaminant Level (ppm)	pH Level	Turbidity (NTU)	Dissolved Oxygen (mg/L)	Nitrate Level (mg/L)
Year								
2000	Mexico	North	Spring	3.97	8.07	1.07	6.37	43.53
2000	India	Central	Pond	2.34	7.53	1.48	4.36	39.79
2000	Nigeria	East	Lake	5.52	8.24	2.83	6.44	40.68
2000	Brazil	East	Well	3.37	6.25	4.41	9.36	35.97
2000	Ethiopia	Central	River	8.39	6.87	2.00	7.17	42.35

5 rows x 23 columns

```
[29]: water_pollution.tail(5)

[29]:
```

	Country	Region	Water Source Type	Contaminant Level (ppm)	pH Level	Turbidity (NTU)	Dissolved Oxygen (mg/L)	Nitrate Level (mg/L)
Year								
2024	Ethiopia	West	Lake	6.35	6.83	4.17	4.02	28.52
2024	China	East	Spring	1.05	6.68	4.31	3.63	46.41
2024	Indonesia	West	Tap	4.62	8.12	0.44	3.01	19.16
2024	Mexico	East	Well	4.20	7.41	1.05	4.11	46.58

- 5: Going back to our dataset, we can see the Country column. We can create sub-dataframe for

each unique country. I first wanted to see what are my unique strings in the Country column. Then I will convert this to a list.

```
[117]: counties_obs = water_pollution["Country"].unique().tolist()
print(counties_obs)

['Mexico', 'India', 'Nigeria', 'Brazil', 'Ethiopia', 'Pakistan', 'Indonesia', 'Bangladesh', 'USA', 'China']
```

7: We can subdivide the water_pollution DF into their respective countries and analyze each one one separately. Since we have 10 different countries, we opt to use a for loop that takes our Country col in water_pollution using the .unique() method. For each 'country' in that list, the code select where "Country" equals that 'country'. The key will be the country's name and its value will be its dataframe.

```
countries_dfs = {country:water_pollution[water_pollution["Country"] == country] for country in water_pollution["Country"].unique()}
```

Key Functions of describe()

- **Numeric Data:** By default, describe() returns statistics for numeric columns, including:
 - count: Number of non-missing values
 - mean: Average value
 - std: Standard deviation
 - min: Minimum value
 - 25%, 50%, 75%: Percentile values (quartiles)
 - max: Maximum value 2 5 6
- **Non-Numeric Data:** For object, categorical, or datetime columns, describe() provides:
 - count: Number of non-missing values
 - unique: Number of unique values
 - top: Most frequent value
 - freq: Frequency of the most frequent value 3 6

8: We can use the describe() to generate descriptive statistics that summarizes the central tendency, dispersion, and shape of a dataset's distribution, excluding NaN values.

8a: We can see the difference in output when applying it to our "pH Level" and "Water Source Type" col.

```
[24]: Central_region["Water Source Type"].describe()
```

```
[24]: count      611  
      unique      6  
      top      Pond  
      freq      107  
      Name: Water Source Type, dtype: object
```

```
[26]: Central_region["pH Level"].describe()
```

```
[26]: count      611.000000  
      mean       7.209509  
      std       0.727203  
      min       6.000000  
      25%       6.540000  
      50%       7.200000  
      75%       7.800000  
      max       8.500000  
      Name: pH Level, dtype: float64
```

```
[36]: len(Central_region["pH Level"])
```

```
[36]: 611
```

rows are in that col.

8b: By applying the len(), we can see how many

```
[38]: Central_region["Water Source Type"].value_counts()
```

```
[38]: Pond      107  
      Tap      102  
      Spring  102  
      Well    101  
      Lake    100  
      River    99  
      Name: Water Source Type, dtype: int64
```

you want to count the frequency of unique values in a Series (such as a DataFrame column). Here we applied the method to see the frequency number of the "Water Source Type" col in the Central_region DF.

8c: We can also calculate

the .median(), .std(), .min() and max().

8d: .nunique() returns the number of unique values in a Series or DF along a specified axis.

The .unique returns a NumPy array of the unique values, in the order they appear.

9: use the value_counts() method in pandas when

```
[50]: Central_region["Temperature (°C)"].nlargest()

[50]: Year
2003    39.98
2006    39.94
2010    39.78
2013    39.61
2024    39.60
Name: Temperature (°C), dtype: float64

[52]: Central_region["Temperature (°C)"].nsmallest()

[52]: Year
2016     0.11
2010     0.16
2011     0.20
2023     0.32
2000     0.33
Name: Temperature (°C), dtype: float64
```

9a: The "normalized" parameter gives us the relative frequency. We can also order the count numbers by using the "ascending" parameter.

10: The nlargest() and nsmallest() method gives us the top n rows of the highest or lowest values. These methods are optimized for speed and memory efficiency. This is a partial sorting algorithms. The default is 5 rows.

```
Name: Temperature (°C), dtype: float64

[20]: #Use bitwise method to get Mexico row only dataframe
Mexico_row = Central_region[Central_region["Country"] == "Mexico"]

[21]: Mexico_row["Temperature (°C)"].nlargest()

[21]: Year
2000    39.39
2007    39.21
2001    39.07
2010    38.88
2005    38.67
Name: Temperature (°C), dtype: float64
```

```
#Creating a conversion function

c_to_f = lambda c: (c* 9/5) + 32

for i in regionsDF :
    i["Temperature"] = i["Temperature"].apply(c_to_f)
```

11: We created a Mexico only dataframe and looked at the highest 5 recorded temperature.

Healthcare Access Index (0- 100)	Urbanization Rate (%)	Sanitation Coverage (% of Population)	Rainfall (mm per year)	Temperature
17.13	49.12	51.28	884	182.8148
10.17	30.41	44.62	583	210.9380
87.84	32.36	40.49	430	123.5876
94.95	72.23	42.97	1379	154.8860
87.33	32.33	56.30	1677	113.9648
...
43.24	87.39	74.88	2395	95.4320
93.18	65.79	68.85	2530	104.0828
80.70	76.62	46.07	788	159.5840
44.43	66.10	73.73	345	192.8264
51.37	78.47	55.32	1475	174.6176

12: We can convert our temperature cols unit from C to F.

```

#Created dict with key:value
region_dict = {
    'NR': North_region,
    'CR': Central_region,
    'ER': East_region,
    'SR': South_region,
    'WR': West_region
}

#Create empty dict
waterTypeFiltered = {
    for i, j in region_dict.items():
        for k in numWaterType:
            key = f"{i}_{k}"
            waterTypeFiltered[key] = j[j["Water Source Type"] == k]
}

```

13: We can now create sub DF based on water

```

for var_name, var_val in (globals().items()):
    if isinstance(var_val, pd.DataFrame):
        print(var_name)

water_pollution
_6
_7
North_region
Central_region
East_region
South_region
West_region
_10
_11
Mexico_row
_51
_59
_75
NR_Lake
NR_Spring
NR_Well
NR_Tap
NR_River
NR_Pond
CR_Lake
CR_Spring
CR_Well
CR_Tap
CR_River
CR_Pond
ER_Lake
ER_Spring
ER_Well
ER_Tap
ER_River
ER_Pond
j

```

source types

```

[97]: North_region["Country"].unique()

[97]: array(['Mexico', 'India', 'Ethiopia', 'Bangladesh', 'Pakistan',
            'Indonesia', 'Brazil', 'China', 'USA', 'Nigeria'], dtype=object)

```

14: Now we can do basic indepth investigation on each region. The first region of focus will be the central region. I used the unique() on the

ANALYSIS OF INDIA

1: We will create a northern India region along with other regions.

```
india_NR = india_df[india_df["Region"] == "North"]
```

india_NR

	Country	Region	Water Source Type	Contaminant Level (ppm)	pH Level	Turbidity (NTU)	Dissolved Oxygen (mg/L)	Nitrate Level (mg/L)	Lead Concentration (µg/L)	Bacteria Count (CFU/mL)
Year										
2000	India	North	Tap	0.28	7.80	1.45	5.07	25.38	3.22	783
2000	India	North	Spring	6.51	7.35	3.09	7.13	38.33	14.12	3808
2000	India	North	Lake	1.23	7.88	1.03	8.82	28.19	1.01	4540
2001	India	North	Pond	8.90	6.20	4.42	5.53	46.08	0.93	3892
2001	India	North	Tap	6.08	7.98	2.09	6.48	21.54	11.57	807
2002	India	North	Well	7.73	7.59	1.26	9.24	42.78	2.79	1776
2003	India	North	Lake	8.41	6.53	0.44	3.63	38.83	9.53	2954
2003	India	North	Pond	4.02	7.92	1.03	3.14	43.48	5.42	830
2004	India	North	Well	4.97	6.91	0.68	9.32	39.88	14.10	3721

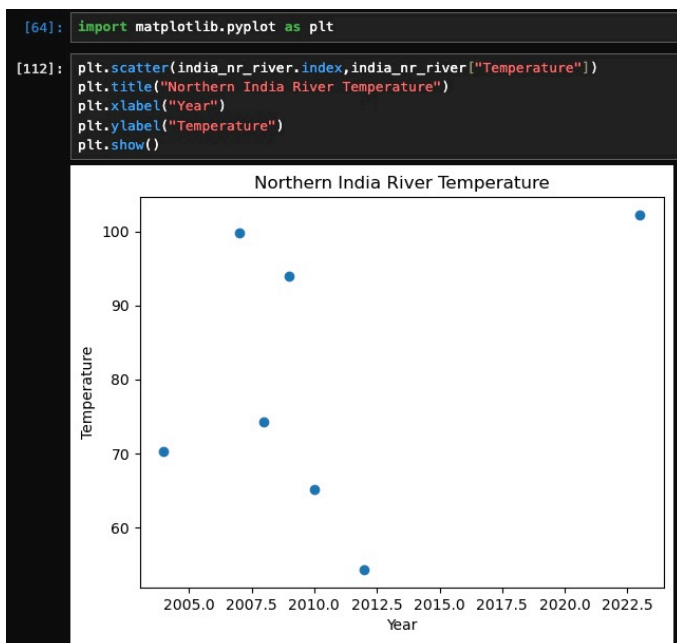
- From the options of water source type, Northern India contains an abundance of rivers and groundwater wells.

WaterAid. (n.d.). Groundwater quality information North India. WASH Matters. <https://washmatters.wateraid.org/publications/groundwater-quality-information-north-india>

Based on these information we will focus on these two water source type

```
[60]: india_nr_river = india_NR[india_NR["Water Source Type"] == "River"]
      india_nr_well = india_NR[india_NR["Water Source Type"] == "Well"]
```

- The first thing I did was observe to see if there is anything meaningful for temperature vs year

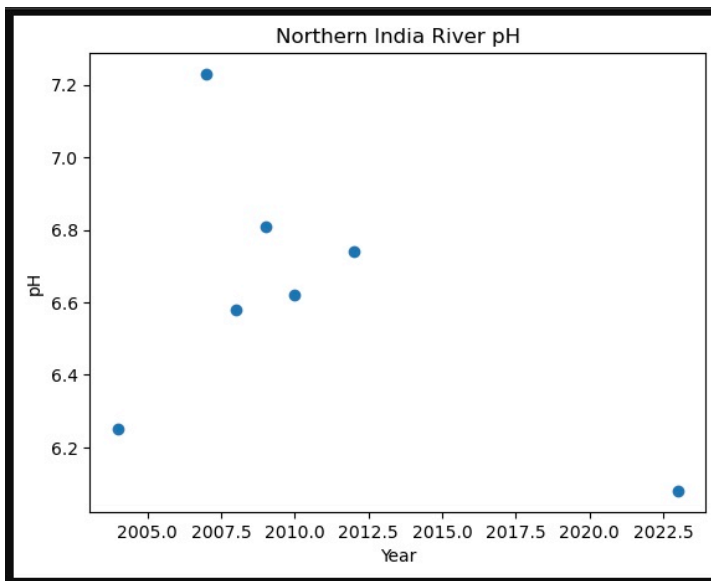


```
[116]: india_nr_river["Temperature"].describe()
```

```
[116]: count      7.000000
      mean       79.985429
      std       18.656252
      min       54.302000
      25%       67.685000
      50%       74.282000
      75%       96.854000
      max      102.236000
      Name: Temperature, dtype: float64
```

- Nothing meaningful was extracted from this graph. Temperature fluctuates with the highest temperature recorded was 102.2F in 2023 and the lowest being 54.3F in 2012.

- The next thing we can observe would be the pH level vs Year



```
[120]: india_nr_river["pH Level"].describe()

[120]: count      7.000000
      mean      6.615714
      std       0.376955
      min       6.080000
      25%       6.415000
      50%       6.620000
      75%       6.775000
      max       7.230000
      Name: pH Level, dtype: float64
```

- From this graph it is difficult to extract any information. The basic stat analysis tells us that the lowest pH recorded was in 2023, which was 6.0. The highest recorded pH was 7.2 on 2007
- It may be difficult to extract anything meaningful due to the numbers of rows. We can combine both the river and well since they are both freshwater.